



Exercise 5.3: Using ConfigMaps Configure Ambassador Containers

In an earlier lab we added a second Ambassador container to handle logging. Now that we have learned about using ConfigMaps and attaching storage we will use configure our basic pod.

1. Review the YAML for our earlier simple pod. Recall that we added an Ambassador style logging container to the pod but had not fully configured the logging.

```
student@master:~$ cat basic.yaml
```

```
1 <output_omitted>
2   containers:
3   - name: webcont
4     image: nginx
5     ports:
6     - containerPort: 80
7   - name: fdlogger
8     image: fluent/fluentd
```

2. Let us begin by adding shared storage to each container. We will use the hostPath storage class to provide the PV and PVC. First we create the directory.

```
student@master:~$ sudo mkdir /tmp/weblog
```

3. Now we create a new PV to use that directory for the hostPath storage class. We will use the storageClassName of manual so that only PVCs which use that name will bind the resource.

```
student@master:~$ vim weblog-pv.yaml
```

YAML

weblog-pv.yaml

```
1 kind: PersistentVolume
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-volume
5   labels:
6     type: local
7 spec:
8   storageClassName: manual
9   capacity:
10    storage: 100Mi
11  accessModes:
12    - ReadWriteOnce
13  hostPath:
14    path: "/tmp/weblog"
```

4. Create and verify the new PV exists and shows an Available status.

```
student@master:~$ kubectl create -f weblog-pv.yaml
```

```
1 persistentvolume/weblog-pv-volume created
```

```
student@master:~$ kubectl get pv weblog-pv-volume
```

```

1 NAME          CAPACITY ACCESS MODES RECLAIM POLICY
2   STATUS      CLAIM   STORAGECLASS  REASON   AGE
3
4 weblog-pv-volume 100Mi    RWO             Retain
5   Available                manual              21s

```

5. Next we will create a PVC to use the PV we just created.

```
student@master:~$ vim weblog-pvc.yaml
```

YAML

weblog-pvc.yaml

```

1 kind: PersistentVolumeClaim
2 apiVersion: v1
3 metadata:
4   name: weblog-pv-claim
5 spec:
6   storageClassName: manual
7   accessModes:
8     - ReadWriteOnce
9   resources:
10    requests:
11      storage: 100Mi

```

6. Create the PVC and verify it shows as Bound to the the PV we previously created.

```
student@master:~$ kubectl create -f weblog-pvc.yaml
```

```
1 persistentvolumeclaim/weblog-pv-claim created
```

```
student@master:~$ kubectl get pvc weblog-pv-claim
```

```

1 NAME          STATUS  VOLUME          CAPACITY ACCESS MODES
2   STORAGECLASS  AGE
3 weblog-pv-claim Bound    weblog-pv-volume 100Mi    RWO
4   manual              79s

```

7. We are ready to add the storage to our pod. We will edit three sections. The first will declare the storage to the pod in general, then two more sections which tell each container where to make the volume available.

```
student@master:~$ vim basic.yaml
```

YAML

basic.yaml

```

1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: basicpod
5   labels:
6     type: webserver
7 spec:
8   volumes:                                     #<-- Add three lines, same depth as containers
9     - name: weblog-pv-storage
10       persistentVolumeClaim:
11         claimName: weblog-pv-claim
12   containers:

```

YAML

```

13   - name: webcont
14     image: nginx
15     ports:
16     - containerPort: 80
17     volumeMounts:                                #<-- Add three lines, same depth as ports
18       - mountPath: "/var/log/nginx/"
19         name: weblog-pv-storage                    # Must match volume name above
20   - name: fdlogger
21     image: fluent/fluentd
22     volumeMounts:                                #<-- Add three lines, same depth as image:
23       - mountPath: "/var/log"
24         name: weblog-pv-storage                    # Must match volume name above

```

8. At this point we can create the pod again. When we create a shell we will find that the `access.log` for **nginx** is no longer a symbolic link pointing to `stdout` it is a writable, zero length file. Leave a **tailf** of the log file running.

```
student@master:~$ kubectl create -f basic.yaml
```

```
1 pod/basicpod created
```

```
student@master:~$ kubectl exec -c webcont -it basicpod -- /bin/bash
```



On Container

```
root@basicpod:/# ls -l /var/log/nginx/access.log
```

```
1 -rw-r--r-- 1 root root 0 Oct 18 16:12 /var/log/nginx/access.log
```

```
root@basicpod:/# tail -f /var/log/nginx/access.log
```

9. Open a second connection to your master node. We will use the pod IP as we have not yet configured a service to expose the pod.

```
student@master:~$ kubectl get pods -o wide
```

```

1 NAME          READY STATUS    RESTARTS  AGE    IP             NODE
2 Nominated node
3 basicpod 2/2   Running  0          3m26s  192.168.213.181 master
4 <none>

```

10. Use **curl** to view the welcome page of the webserver. When the command completes you should see a new entry added to the log. Right after the GET we see a 200 response indicating success. You can use **ctrl-c** and **exit** to return to the host shell prompt.

```
student@master:~$ curl http://192.168.213.181
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <output_omitted>

```



On Container

```
192.168.32.128 - - [18/Oct/2018:16:16:21 +0000] "GET / HTTP/1.1" 200 612 "-" "curl/7.47.0" "-"
```

11. Now that we know the webcont container is writing to the PV we will configure the logger to use that directory as a source. For greater flexibility we will configure **fluentd** using a configMap.

Fluentd has many options for input and output of data. We will read from a file of the webcont container and write to standard out of the fdlogger container. The details of the data settings can be found in **fluentd** documentation here: <https://docs.fluentd.org/v1.0/categories/config-file>

```
student@master:~$ vim weblog-configmap.yaml
```



weblog-configmap.yaml

```
1 apiVersion: v1
2 kind: ConfigMap
3 metadata:
4   name: fluentd-config
5   namespace: default
6 data:
7   fluentd.conf: |
8     <source>
9       @type tail
10      format none
11      path /var/log/access.log
12      tag count.format1
13    </source>
14
15    <match *.**>
16      @type stdout
17      id stdout_output
18    </match>
```

12. Create the new configMap.

```
student@master:~$ kubectl create -f weblog-configmap.yaml
```

```
1 configmap/fluentd-config created
```

13. View the logs for both containers in the basicpod. You should see some startup information, but not the HTTP traffic.

```
student@master:~$ kubectl logs basicpod webcont
```

```
1 /docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
2 /docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
3 /docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
4 10-listen-on-ipv6-by-default.sh: Getting the checksum of /etc/nginx/conf.d/default.conf
5 10-listen-on-ipv6-by-default.sh: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
6 /docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
7 /docker-entrypoint.sh: Configuration complete; ready for start up
```

```
student@master:~$ kubectl logs basicpod fdlogger
```

```
1 2020-09-02 19:32:59 +0000 [info]: reading config file path="/etc/fluentd-config/fluentd.conf"
2 2020-09-02 19:32:59 +0000 [info]: starting fluentd-0.12.29
3 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-config-placeholders' version '0.4.0'
4 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-plaintextformatter' version '0.2.6'
5
6 <output_omitted>
```

```

7
8 <source>
9   @type tail
10  format none
11  path /var/log/access.log
12 <output_omitted>

```

14. Now we will edit the pod yaml file so that the **fluentd** container will mount the configmap as a volume and reference the variables inside the config file. You will add three areas, the volume declaration to the pod, the `env` parameter and the mounting of the volume to the fluentd container

```
student@master:~$ vim basic.yaml
```

YA
ML

basic.yaml

```

1  ....
2  volumes:
3    - name: weblog-pv-storage
4      persistentVolumeClaim:
5        claimName: weblog-pv-claim
6    - name: log-config                    #<-- This and two lines following
7      configMap:
8        name: fluentd-config             # Must match existing configMap
9  ....
10  image: fluent/fluentd
11  env:                                   #<-- This and two lines following
12    - name: FLUENTD_OPT
13      value: -c /etc/fluentd-config/fluentd.conf
14  ....
15  volumeMounts:
16    - mountPath: "/var/log"
17      name: weblog-pv-storage
18    - name: log-config                    #<-- This and next line
19      mountPath: "/etc/fluentd-config"

```

15. At this point we can delete and re-create the pod, which would cause the configmap to be used by the new pod, among other changes.

```
student@master:~$ kubectl delete pod basicpod
```

```
1 pod "basicpod" deleted
```

```
student@master:~$ kubectl create -f basic.yaml
```

```
1 pod/basicpod created
```

```
student@master:~$ kubectl get pod basicpod -o wide
```

```

1 NAME          READY   STATUS    RESTARTS   AGE   IP              NODE   NOMINATED....
2 basicpod      2/2     Running   0           8s    192.168.171.122 worker   <none>   ....

```

16. Use **curl** a few times to look at the default page served by basicpod

```
student@master:~$ curl http://192.168.171.122
```

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6   body {
7 <output_omitted>

```

17. Look at the logs for both containers. In addition to the standard startup information, you should also see the HTTP requests from the curl commands you just used at the end of the fdlogger output.

```
student@master:~$ kubectl logs basicpod webcont
```

```

1 /docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
2 /docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
3 /docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
4 10-listen-on-ipv6-by-default.sh: Getting the checksum of /etc/nginx/conf.d/default.conf
5 10-listen-on-ipv6-by-default.sh: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
6 /docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
7 /docker-entrypoint.sh: Configuration complete; ready for start up

```

```
student@master:~$ kubectl logs basicpod fdlogger
```

```

1 2020-09-02 19:32:59 +0000 [info]: reading config file path="/etc/fluentd-config/fluentd.conf"
2 2020-09-02 19:32:59 +0000 [info]: starting fluentd-0.12.29
3 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-config-placeholders' version '0.4.0'
4 2020-09-02 19:32:59 +0000 [info]: gem 'fluent-mixin-plaintextformatter' version '0.2.6'
5
6 <output_omitted>
7
8 <source>
9   @type tail
10  format none
11  path /var/log/access.log
12
13 <output_omitted>
14
15 2020-09-02 19:47:38 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:38 +0000] \"GET / HTTP/1.1\""}
16 2020-09-02 19:47:41 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:41 +0000] \"GET / HTTP/1.1\""}
17 2020-09-02 19:47:47 +0000 count.format1: {"message":"192.168.219.64 - - [02/Sep/2020:19:47:47 +0000] \"GET / HTTP/1.1\""}

```